

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
df = pd.read_csv("ecommerce.csv")
```

```
df.head()
```

	Email \
0	mstephenson@fernandez.com
1	hduke@hotmail.com
2	pallen@yahoo.com
3	riverarebecca@gmail.com
4	mstephens@davidson-herman.com

	Address	Avatar
0	835 Frank Tunnel\nWrightmouth, MI 82180-9605	Violet
1	4547 Archer Common\nDiazchester, CA 06566-8576	DarkGreen
2	24645 Valerie Unions Suite 582\nCobbborough, D...	Bisque
3	1414 David Throughway\nPort Jason, OH 22070-1220	SaddleBrown
4	14023 Rodriguez Passage\nPort Jacobville, PR 3...	MediumAquaMarine

	Avg. Session Length	Time on App	Time on Website	Length of Membership \
0	34.497268	12.655651	39.577668	4.082621
1	31.926272	11.109461	37.268959	2.664034
2	33.000915	11.330278	37.110597	4.104543
3	34.305557	13.717514	36.721283	3.120179
4	33.330673	12.795189	37.536653	4.446308

	Yearly Amount Spent
0	587.951054
1	392.204933
2	487.547505
3	581.852344
4	599.406092

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 500 entries, 0 to 499
Data columns (total 8 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Email                                500 non-null    object
1   Address                             500 non-null    object
2   Avatar                              500 non-null    object
3   Avg. Session Length                 500 non-null    float64
4   Time on App                         500 non-null    float64
5   Time on Website                     500 non-null    float64
6   Length of Membership                500 non-null    float64
7   Yearly Amount Spent                 500 non-null    float64
dtypes: float64(5), object(3)
memory usage: 31.4+ KB
```

```
df.describe()
```


	Avg. Session Length	Time on App	Time on Website \
count	500.000000	500.000000	500.000000
mean	33.053194	12.052488	37.060445
std	0.992563	0.994216	1.010489
min	29.532429	8.508152	33.913847
25%	32.341822	11.388153	36.349257
50%	33.082008	11.983231	37.069367
75%	33.711985	12.753850	37.716432
max	36.139662	15.126994	40.005182

	Length of Membership	Yearly Amount Spent
count	500.000000	500.000000
mean	3.533462	499.314038
std	0.999278	79.314782
min	0.269901	256.670582
25%	2.930450	445.038277
50%	3.533975	498.887875
75%	4.126502	549.313828
max	6.922689	765.518462

```
#EDA
```


```
sns.jointplot(x="Time on Website" , y="Yearly Amount Spent", data =df,
alpha = 0.5)
```

```
<seaborn.axisgrid.JointGrid at 0x24dfb572cf0>
```




```
sns.jointplot(x= "Time on App" , y="Yearly Amount Spent", data = df,  
alpha = 0.5)
```

```
<seaborn.axisgrid.JointGrid at 0x24dfb816350>
```




```
sns.pairplot(df,kind="scatter",plot_kws={"alpha" : 0.4})  
<seaborn.axisgrid.PairGrid at 0x24dfb68f4d0>
```



```
sns.lmplot(x='Length of Membership',
           y='Yearly Amount Spent',
           data = df,
           scatter_kws = {'alpha':0.3})
```

<seaborn.axisgrid.FacetGrid at 0x24dfb68e510>



```
from sklearn.model_selection import train_test_split

X = df[['Time on App', 'Time on Website', 'Avg. Session Length', 'Length
of Membership']]
y= df['Yearly Amount Spent']
X
```

	Time on App	Time on Website	Avg. Session Length	Length of Membership
0	12.655651	39.577668	34.497268	
4.082621				
1	11.109461	37.268959	31.926272	
2.664034				
2	11.330278	37.110597	33.000915	
4.104543				
3	13.717514	36.721283	34.305557	
3.120179				
4	12.795189	37.536653	33.330673	
4.446308				
...	
...				

495	13.566160	36.417985	33.237660
3.746573			
496	11.695736	37.190268	34.702529
3.576526			
497	11.499409	38.332576	32.646777
4.958264			
498	12.391423	36.840086	33.322501
2.336485			
499	12.418808	35.771016	33.715981
2.735160			

[500 rows x 4 columns]

y

0	587.951054
1	392.204933
2	487.547505
3	581.852344
4	599.406092
	...
495	573.847438
496	529.049004
497	551.620145
498	456.469510
499	497.778642

Name: Yearly Amount Spent, Length: 500, dtype: float64

X_train, X_test, y_train, y_test = train_test_split(X,y,test_size = 0.3, random_state = 42)

X_train

	Time on App	Time on Website	Avg. Session Length	Length of Membership
5	12.026925	34.476878	33.871038	
5.493507				
116	12.011022	36.701052	33.925795	
2.753424				
45	12.170525	39.131097	34.555768	
3.663105				
16	11.733862	34.894093	32.125387	
3.136133				
462	11.233415	37.211153	33.503810	
2.320550				
..	
...				
106	12.190474	36.152462	32.291756	
3.781823				
270	12.956277	38.655095	34.006489	

```
3.275734
348    10.886921    34.897828    31.812483
3.128639
435    14.132893    37.023479    32.259973
3.762070
102    11.448902    37.580190    32.425697
2.586968
```

```
[350 rows x 4 columns]
```

```
X_test
```

```
      Time on App  Time on Website  Avg. Session Length  Length of
Membership
361    10.347877    39.045156    32.077590
3.434560
73     12.817113    37.031539    32.808698
3.851579
374    10.101632    38.043453    31.447446
4.238296
155    13.457725    37.238806    32.449522
2.941411
104    10.994224    38.074452    31.389585
3.428860
...      ...      ...      ...
...
266    11.777772    37.979827    34.555283
3.784273
23     11.657576    36.772604    32.903251
3.919302
222    11.109456    38.585855    34.334865
3.892891
261    13.041245    36.655208    32.550527
3.456234
426    13.271475    37.239847    31.425227
4.022103
```

```
[150 rows x 4 columns]
```

```
y_train
```

```
5      637.102448
116    479.231093
45     549.860590
16     457.847696
462    397.420584
...
106    494.551861
270    540.995739
348    392.810345
```



```
435    571.216005
102    420.737673
Name: Yearly Amount Spent, Length: 350, dtype: float64
```

```
#training the model
```

```
from sklearn.linear_model import LinearRegression
```

```
lm=LinearRegression()
```

```
lm.fit(X_train,y_train)
```

```
LinearRegression()
```

```
lm.coef_
```

```
array([38.59713548,  0.45914788, 25.72425621, 61.67473243])
```

```
cdf=pd.DataFrame(lm.coef_, X.columns, columns=['Coef'])
```

```
cdf
```

	Coef
Time on App	38.597135
Time on Website	0.459148
Avg. Session Length	25.724256
Length of Membership	61.674732

```
#predictions
```


```
predictions=lm.predict(X_test)
```

```
predictions
```

```
array([403.66993069, 542.57756289, 427.06591658, 502.02460425,
       410.12143559, 569.93442508, 531.93431341, 506.29650969,
       408.71870658, 473.97737105, 441.46912726, 425.33703059,
       425.1297229 , 527.61676714, 431.45684016, 424.0769184 ,
       575.76543296, 484.89856554, 458.35936863, 481.96502182,
       502.32441491, 513.63783554, 507.58877002, 646.57464283,
       450.24372141, 496.27043415, 556.40457807, 554.95630839,
       399.64237199, 325.84623136, 532.89783259, 478.12238702,
       501.05701845, 305.97335848, 505.77244448, 483.79591969,
       518.8331528 , 438.18241857, 456.71094234, 471.04609461,
       494.44008972, 445.31155755, 508.78802753, 501.04594193,
       488.83499673, 535.38079541, 595.20129802, 514.04714872,
       280.76758312, 433.10112367, 421.70823427, 481.23640152,
       584.71372272, 608.7748096 , 563.98513427, 494.72804869,
       394.52133407, 456.4197529 , 573.08767515, 499.6984241 ,
       512.83277025, 392.12434043, 480.05057697, 481.54520299,
       475.1117359 , 546.2717533 , 430.85039085, 602.16082001,
       422.3695128 , 493.57280186, 528.74970313, 581.49002635,
       620.19139276, 512.56880298, 411.76623862, 498.47637494,
       461.51337557, 446.41371051, 448.07229961, 535.44710412,
```

```
599.45225302, 619.33717662, 494.15919062, 671.99976398,  
532.46469814, 438.90606319, 515.04975242, 546.7821954 ,  
331.94282076, 510.51987447, 536.57891032, 500.19533618,  
376.92345776, 573.73961388, 479.68031607, 588.61435483,  
485.69922203, 456.40200844, 399.25197845, 451.5098931 ,  
519.40693826, 434.71194217, 596.13049586, 487.91791966,  
407.46691799, 524.16812757, 504.12982787, 452.11540623,  
524.21791295, 457.59311643, 444.19371592, 457.80432916,  
448.76590761, 438.31789012, 677.04967982, 566.09639245,  
651.93616661, 381.08127926, 577.5577254 , 578.35797052,  
518.61431291, 538.94532336, 377.4301223 , 663.30814872,  
523.83158824, 456.86065622, 446.07594402, 388.55038282,  
521.03242183, 431.94999241, 460.08016327, 426.31959507,  
433.30417088, 634.89577554, 462.41086078, 460.71673829,  
512.49535288, 703.83033889, 411.84238624, 551.54681408,  
553.33669558, 409.68202123, 423.34491341, 509.66438623,  
509.88865178, 543.67591782, 504.31300469, 519.18802223,  
520.03155195, 535.13855037])
```

```
sns.scatterplot(x=predictions,y=y_test)  
plt.xlabel('predictions')  
plt.title("Evaluation of our LM Model")  
Text(0.5, 1.0, 'Evaluation of our LM Model')
```



```
from sklearn.metrics import mean_squared_error, mean_absolute_error
import math

print("Mean Absolute Error: ",
mean_absolute_error(y_test,predictions))
print("Mean Squared Error: ", mean_squared_error(y_test,predictions))
print("RMSE: ", math.sqrt(mean_squared_error(y_test,predictions)))
```

```
Mean Absolute Error:  8.4260916414321
Mean Squared Error:  103.91554136503328
RMSE:  10.193897260863153
```


```
#residual analysis
residuals = y_test-predictions
residuals
```

```
361    -2.636795
73     -7.800375
374    -8.463174
155     1.953775
104    -0.051825
...
266    10.327176
23     15.027984
```

```
222    -16.778237
261    -6.021734
426    -4.371832
Name: Yearly Amount Spent, Length: 150, dtype: float64
```

```
#test normality of residuals
sns.displot(residuals, bins=30,kde=True)

<seaborn.axisgrid.FacetGrid at 0x24d85413750>
```



```
import pylab
import scipy.stats as stats
stats.probplot(residuals,dist='norm',plot=pylab)
pylab.show()
```

